

Fnu Suya

✉ suya@umd.edu • fsuya.org • [in suya](https://in.suya) • [suyafnu](https://twitter.com/suyafnu)
🔗 [suyeeecav](#) • Updated on September 24, 2026

Professional Appointments

Tenure-track Assistant Professor 08/2024 – Present
Min H. Kao Department of EECS, University of Tennessee, Knoxville
MC2 Postdoctoral Fellow 08/2023 – 07/2024
Maryland Cybersecurity Center, University of Maryland, College Park

Education

Ph.D. in Computer Science, *University of Virginia* 08/2017 – 07/2023
Thesis: *On the Limits of Data Poisoning Attacks*
Advisor: [Yuan Tian](#) (UCLA), [David Evans](#)
M.Sc. in Computer Science, *University of Virginia* 08/2017 – 12/2022
Advisor: [Yuan Tian](#) (UCLA), [David Evans](#)
Ph.D. in Computer Science, *Arizona State University* 08/2015 – 07/2017
Advisor: [Paolo Papotti](#) (EURECOM), discontinued due to advisor relocation
B.Eng. in Electrical Engineering (Honors), *China Agricultural University* 09/2010 – 07/2014

Research Interests

Trustworthy Machine Learning, Machine Learning for Security

Internship Experience

Research Intern, *Qualcomm*, San Diego (with [Aleksandr Triastcyn](#) on robust FL) 05/2021– 08/2021
Research Intern, *Amazon AWS*, New York City (with [MohamadAli Torkamani](#) on GNNs) 01/2021 – 04/2021
Research Intern, *Bosch AI Center*, Pittsburgh (with [Anit Kumar Sahu](#) on black-box attacks) 06/2020 – 08/2020
Research Intern, *Tsinghua University*, China (with [Bo Bai](#) on wireless communication) 08/2014 – 02/2015

Selected Publications

2026a S. Hoque, J. Zhang, J. Sun, **F. Suya**. “Token Inflation: How Dishonest Providers Can Overcharge for Large Language Model Usage”. In: *NeurIPS 2026*.
2026b H. Rahaman, Q. Xu, Z. Haider, P. Chakraborty, S. Bhunia, **F. Suya**. “(A)iSpy: Parasitic Trojans for Machine Learning Infrastructure”. In: *ACM CCS 2026*.
2026c T. Zhang, **F. Suya**, R. Jha, C. Zhang, V. Shmatikov. “Adversarial Hubness in Multi-Modal Retrieval”. In: *IEEE S&P (Oakland) 2026*.
2026d S. Amgain, D. Lobo, A. Chatterjee, S. Bhunia, **F. Suya**. “HAMLOCK: HArdware-Model Logically Combined attackK”. In: *USENIX Security 2026*.
2024 **F. Suya**, A. Suri, T. Zhang, S. Hong, Y. Tian, D. Evans. “SoK: Pitfalls in Evaluating Black-Box Attacks”. In: *IEEE SaTML 2024*.
2023 **F. Suya**, X. Zhang, Y. Tian, D. Evans. “What Distributions are Robust to Indiscriminate Poisoning Attacks for Linear Learners?”. In: *NeurIPS 2023*.

2022 Y. Tian, **F. Suya**, F. Xu, D. Evans. "Stealthy Backdoors as Compression Artifacts". In: *IEEE TIFS* 2022.

Conference and Journal Papers

Google Scholar ID: [OmLIG8EAAAJ](#)

- 2027 J. Zhang, Y. He, K. Cai, H. Zhao, **F. Suya**, Y. Tian. "RogueMerge: Robust and Unified Attacks against LLM Model Merging". In: *IEEE S&P (Oakland)* 2027.
- 2026a A. Al Nomaan Nafi, H. Rahaman, Z. Haider, T. Mahfuz, **F. Suya**, S. Bhunia, P. Chakraborty. "DAASH: A Meta-Attack Framework for Synthesizing Effective and Stealthy Adversarial Examples". In: *CVPR* 2026.
- 2026b S. Amgain, D. Lobo, A. Chatterjee, S. Bhunia, **F. Suya**. "HAMLOCK: HARDware-Model Logically Combined attackK". In: *USENIX Security* 2026.
- 2026c S. Hoque, J. Zhang, J. Sun, **F. Suya**. "Token Inflation: How Dishonest Providers Can Overcharge for Large Language Model Usage". In: *NeurIPS* 2026.
- 2026d H. Rahaman, Q. Xu, Z. Haider, P. Chakraborty, S. Bhunia, **F. Suya**. "(A)iSpy: Parasitic Trojans for Machine Learning Infrastructure". In: *ACM CCS* 2026.
- 2026e T. Zhang, **F. Suya**, R. Jha, C. Zhang, V. Shmatikov. "Adversarial Hubness in Multi-Modal Retrieval". In: *IEEE S&P (Oakland)* 2026.
- 2024 **F. Suya**, A. Suri, T. Zhang, S. Hong, Y. Tian, D. Evans. "SoK: Pitfalls in Evaluating Black-Box Attacks". In: *IEEE SaTML* 2024.
- 2023a **F. Suya**, X. Zhang, Y. Tian, D. Evans. "What Distributions are Robust to Indiscriminate Poisoning Attacks for Linear Learners?" In: *NeurIPS* 2023.
- 2023b Y. Tian, **F. Suya**, A. Suri, F. Xu, D. Evans. "Manipulating Transfer Learning for Property Inference". In: *CVPR* 2023.
- 2022 Y. Tian, **F. Suya**, F. Xu, D. Evans. "Stealthy Backdoors as Compression Artifacts". In: *IEEE TIFS*.
- 2021 **F. Suya**, S. Mahloujifar, A. Suri, D. Evans, Y. Tian. "Model-Targeted Poisoning Attacks with Provable Convergence". In: *ICML* 2021.
- 2020a **F. Suya**, J. Chi, D. Evans, Y. Tian. "Hybrid Batch attacks: Finding Black-box Adversarial Examples with Limited Queries". In: *USENIX Security* 2020.
- 2020b J. Wang, M. Luo, **F. Suya**, J. Li, Z. Yang, Q. Zheng. "Scalable Attack on Graph Data by Injecting Vicious Nodes". In: *ECML-PKDD* 2020.
- 2019 Y. Chen, M. Zha, N. Zhang, D. Xu, Q. Zhao, X. Feng, K. Yuan, **F. Suya**, Y. Tian, K. Chen, X. Wang, W. Zou. "Demystifying Hidden Privacy Settings in Mobile Apps". In: *IEEE S&P (Oakland)* 2019.
- 2016 **F. Suya**, Y. Shi, B. Bai, W. Chen, J. Zhang, K. B. Letaief, S. Zhou. "Optimal Stochastic Power Control with Compressive CSI Acquisition for Cloud-RAN". In: *GlobalSIP* 2016.

Preprints

- 2026a G. Lakshmanan and **F. Suya**. *Are Commercial Moderation Systems Stable under Semantic-Preserving Stylistic Variation?* Manuscript.
- 2026b A. Toumazatos, G. Syros, Y. Kaya, **F. Suya**, A. Deshpande, T. Dumitras. *Deconstructing Bao: A Study of Feature Impact and Model Robustness in Learned Query Optimization*. Manuscript.
- 2026c K. Cai, K. Zhang, Z. Xiang, J. Zhang, A. Alwan, **F. Suya**, Y. Tian. *Hearing the Whispers: Black-Box Membership Inference Attacks on Finetuned TTS Models*. arXiv preprint arXiv:2609.01723.
- 2026d L. H. Rodoshi, F. Al Rifat, Y. Li, **F. Suya**, F. H. Shezan. *ReDit: Automated Record-Keeping Audit of AI Agent under the EU AI Act*. Manuscript.
- 2026e J. Huang, C. Gong, Y. Pang, **F. Suya**. *Demystifying Prompt Stealing Attacks against Synthetic Video*. Manuscript.
- 2026f S. Kwag, **F. Suya**, T. Dumitras. *Split-Personality Disorders: The Impact of Intra-Sample Concept Drift on Behavioral Malware Analysis in the Wild*. Manuscript.
- 2026g J. Zhang, P. Yu, Z. Lin, K. Cai, **F. Suya**, Y. Tian. *ImageAuditor: Membership Inference Attack against Image-based Retrieval-Augmented Generation*. arXiv preprint arXiv:2606.03354.
- 2026h A. Al Nomaan Nafi, **F. Suya**, S. Bhunia, P. Chakraborty. *LASH: Adaptive Semantic Hybridization for Black-Box Jailbreaking of Large Language Models*. arXiv preprint arXiv:2605.21362.

- 2026i** A. A. Taki, A. Wang, **F. Suya**. *ProxMatch: Feature-Space Proximity for Calibrated Pseudo-Label Propagation in Semi-Supervised Learning*. Manuscript.
- 2026j** X. Xiao, T.-T. Gao, H. You, X. Hu, H. Ping, S. Li, Y. Dong, Y. Zhao, H. Hang, **F. Suya**, T. Liu, P. Zhang, S. Li, S. Qin, W. Yang, X. Zhang, H. Xu, X. Mao, Y. Zhao, N. Kotov. *Large AI Models Need Structural Metrology*. Manuscript.
- 2026k** K. Cai, R. Song, J. Zhang, K. Zhang, P. Bodapati, A. Yu, **F. Suya**, M. Rostami, J. Ma, Y. Tian. *What-If World: A Causal Benchmark for General World Models in Embodied Scenarios*. arXiv preprint arXiv:2605.27589.

Workshop Papers and Posters

- 2023a** A. Kinf, C. Jung, K. Lin, M. Clyburn, **F. Suya**. "HackWrt: Network Traffic-Based Eavesdropping of Handwriting". In: *CPS-IoT Week 2023*.
- 2023b** **F. Suya**, X. Zhang, Y. Tian, D. Evans. "When Can Linear Learners be Robust to Indiscriminate Poisoning Attacks?" In: *ICML AdvML-Frontiers Workshop 2023*.
- 2022** E. Rose, **F. Suya**, D. Evans. "Poisoning Attacks and Subpopulation Susceptibility". In: *VISxAI workshop 2022 (Best Paper Award)*.
- 2018** **F. Suya**, D. Evans, Y. Tian. "Poster: Adversaries Don't Care About Averages: Batch Attacks on Black-Box Classifiers". In: *IEEE S&P (Oakland) 2018*.
- 2017** **F. Suya**, Y. Tian, D. Evans, P. Papotti. "Query-limited Black-box Attacks to Classifiers". In: *N(eur)IPS MLSec Workshop 2017 (Spotlight Presentation)*.

Talks and Presentations

1. An Adversarial Perspective on Machine Learning Pipeline
(Invited Talk) Warren B. Nelms Annual IoT Conference 2024, Gainesville, FL 12/2024
2. What Distributions are Robust to Indiscriminate Poisoning Attacks for Linear Learners?
NeurIPS, New Orleans, LA 12/2023
3. When Can Linear Learners be Robust to Indiscriminate Poisoning Attacks?
ICML Workshop on AdvML-Frontiers, Online 07/2023
4. Model-Targeted Poisoning Attacks with Provable Convergence
ICML, Online 07/2021
5. Hybrid Batch Attacks: Finding Black-box Adversarial Examples with Limited Queries
Usenix Security, Online 08/2020
6. Adversaries Don't Care About Averages: Batch Attacks on Black-Box Classifiers
IEEE S&P (Oakland), San Francisco, CA 05/2018
7. Query-limited Black-box Attacks to Classifiers
N(eur)IPS workshop on MLSec, Long Beach, CA 12/2017
8. Optimal Stochastic Power Control with Compressive CSI Acquisition for Cloud-RAN
GlobalSIP, Washington DC 12/2016

Mentoring Experience

- 2026 - Present, Ony Hoque, PhD student, University of Tennessee, Knoxville (co-advised with Jinyuan Stella Sun)
- 2025 - Present, Sanskar Amgain, PhD student, University of Tennessee, Knoxville
- 2025 - Present, Qipan Xu, MS student, University of Tennessee, Knoxville
- 2024 - 2026, Gomathi Lakshmanan, MS student, University of Tennessee, Knoxville (graduated)

Teaching Experience

- Machine learning (UTK COSC 522), Instructor F2026
- Introduction to Cybersecurity (UTK COSC 366), Instructor S2026

○ Machine learning (UTK COSC 522), Instructor	F2025
○ Adversarial Machine learning (UTK COSC 594), Instructor	S2025
○ Introduction to Cybersecurity (UTK COSC 366), Instructor	F2024
○ Learning Theory (UVA CS 6501-005), TA	S2019
○ Cryptography (UVA CS 6501-009), TA	S2019
○ Game Theory (ASU CSE 556), TA	F2016
○ Introduction to C++ Programming (ASU CSE 100), TA	F2015 – S2016
○ Introduction to Programming Languages (ASU CSE 240), TA	F2015 – S2016

Grants

- National Institutes of Health (NIH) research award (Role: Co-PI; Amount: \$802,713) 2026
- DOE Genesis Mission, U.S. Department of Energy, led by North Carolina State University (Role: UTK Lead PI; UTK Amount: \$287,000) 2026
- Energy & Gateway Internal Award, ORIED, University of Tennessee (Role: PI; Amount: \$60,000) 2026
- Gift from Arista Networks in support of the TRUST-AI Seminar Series (Amount: \$10,000) 2026
- Support for Affiliated Research Teams (StART), UT-Oak Ridge Innovation Institute (Role: PI; Amount: \$200,000; Funding Rate: 6%) 2025 – 2026
- AI TechX Award, AI Tennessee Initiative, University of Tennessee, jointly with Cisco (Role: PI; Amount: \$60,000) 2025
- ORIED Internal Award in support of the TRUST-AI Seminar Series, ORIED, University of Tennessee (Amount: \$10,000) 2025

Honors & Awards

- MC2 Postdoctoral Fellowship: awarded annually to one individual by the Maryland Cybersecurity Center ([MC2](#)) 2023
- Outstanding/Top Reviewer: ICLR [2022](#), NeurIPS [2022](#), [2023](#)
- CS Graduate Research Award, University of Virginia 2018
- CS Department Fellowship, University of Virginia 2017
- NSF Travel Grant 2016
- CIDSE Doctoral Fellowship, Arizona State University 2015
- Outstanding Student Scholarship, China Agricultural University 2011 – 2013

Service

Associate Editor	Journal of Hardware and Systems Security (Springer), 2026 – Present
PC/Reviewer	IEEE S&P, Usenix Security, CCS, NDSS, IEEE EuroS&P, AsiaCCS, ICML, NeurIPS, ICLR, ICCV, CVPR, ECCV, AISTATS, IJCAI, AAAI
Journal Reviewer	Artificial Intelligence, TMLR

Skills

Programing Languages	Python, PyTorch, TensorFlow, Matlab, C, C++ Mongolian (native), Mandarin, English
----------------------	--

References

- Available upon request